



# AMQP/Qpid

**JUG 2011**

**Arnaud Simon**  
**[asimon@redhat.com](mailto:asimon@redhat.com)**

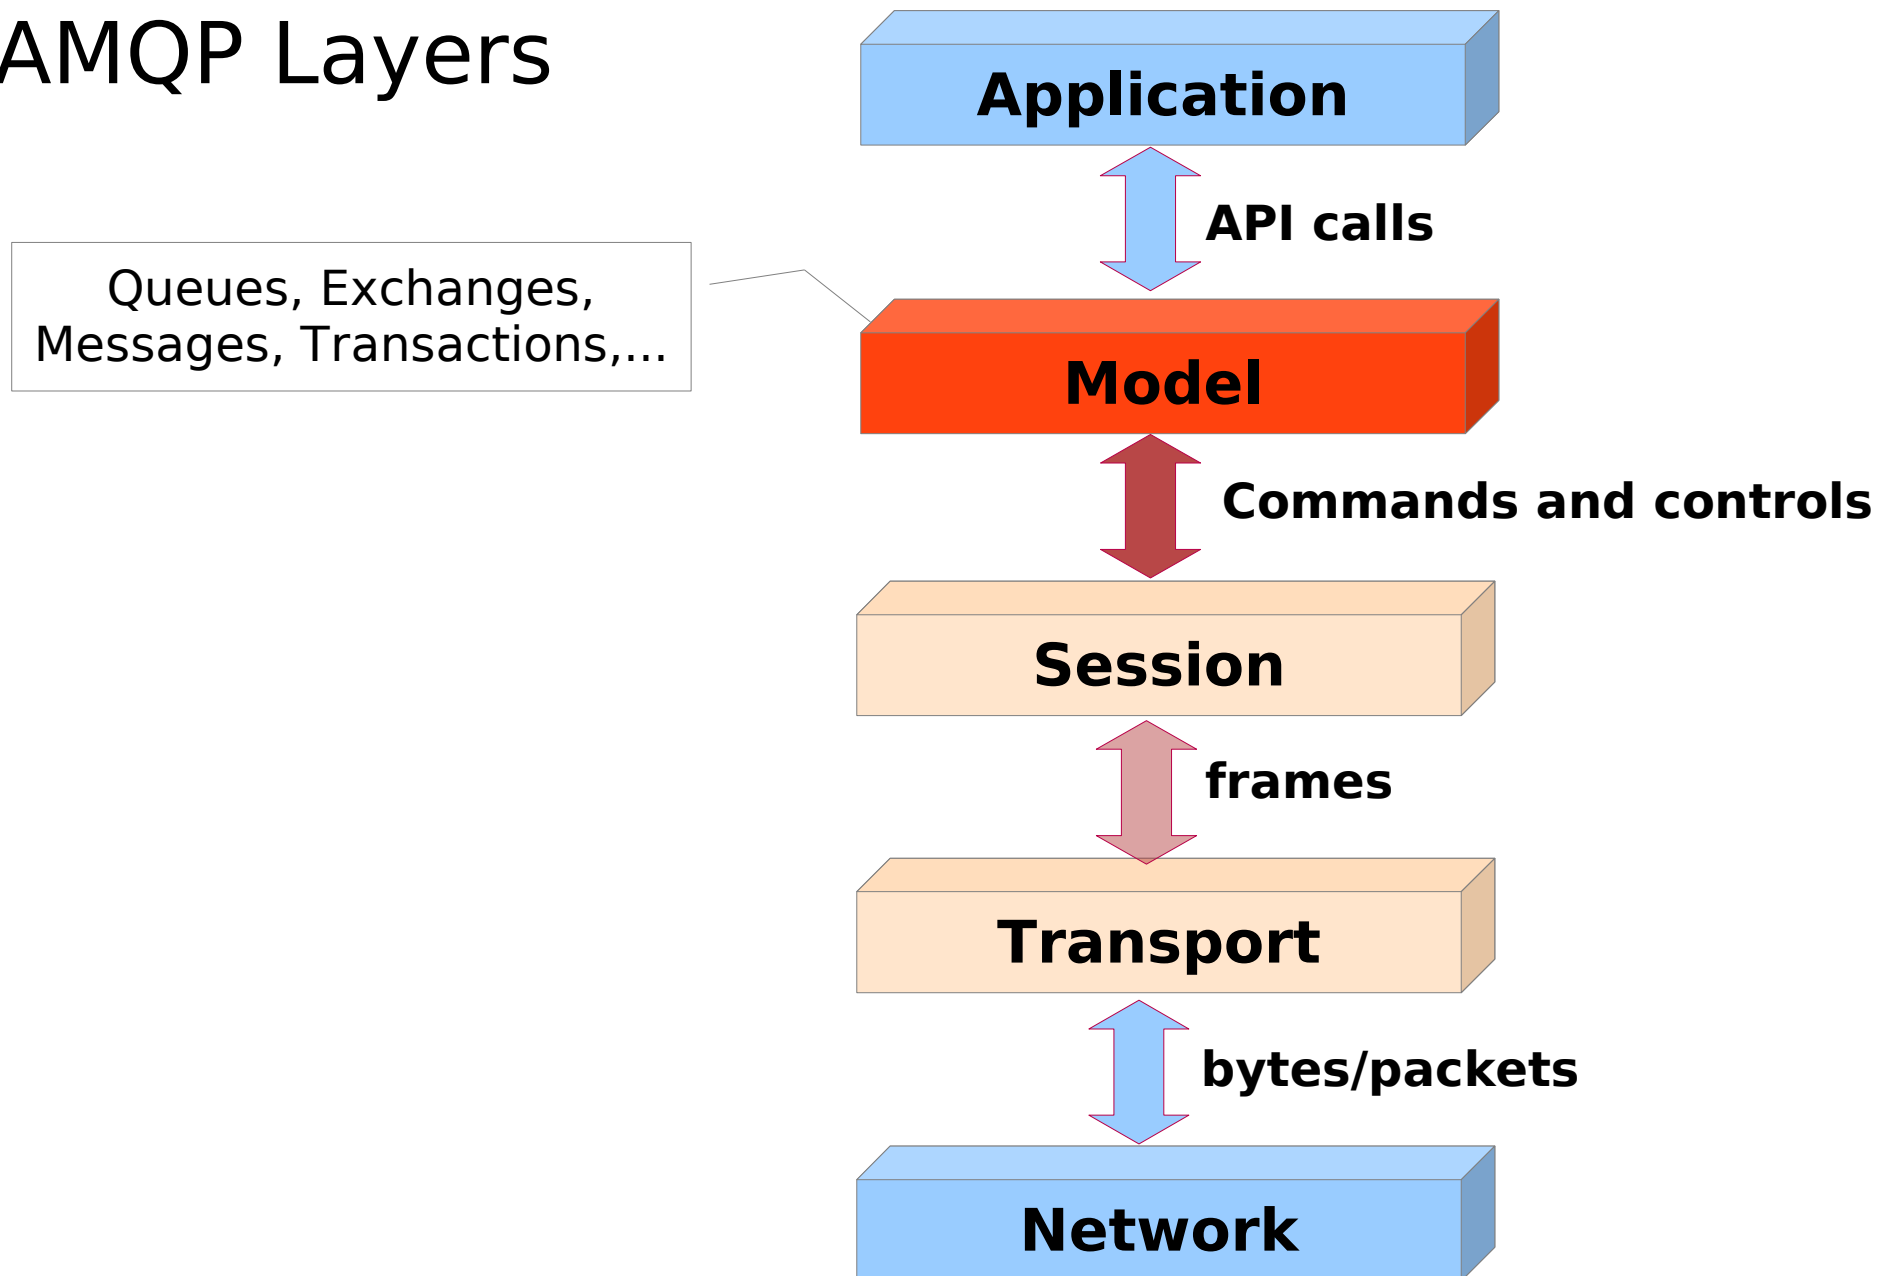
# AMQP: Advanced Message Queuing Protocol

- An open standard for Messaging Middleware
- Pervasive deployment:
  - full interoperability
    - Across platforms, languages, vendors
    - Drop-in compatible with Java JMS
  - Message exchange semantics
  - Network protocol
- Complete solution for business messaging:
  - High performances, Robust, available, Scalable, Secure, Transacted, secure, resilient, ...
- Created by users and technologists

## Who is behind AMQP?

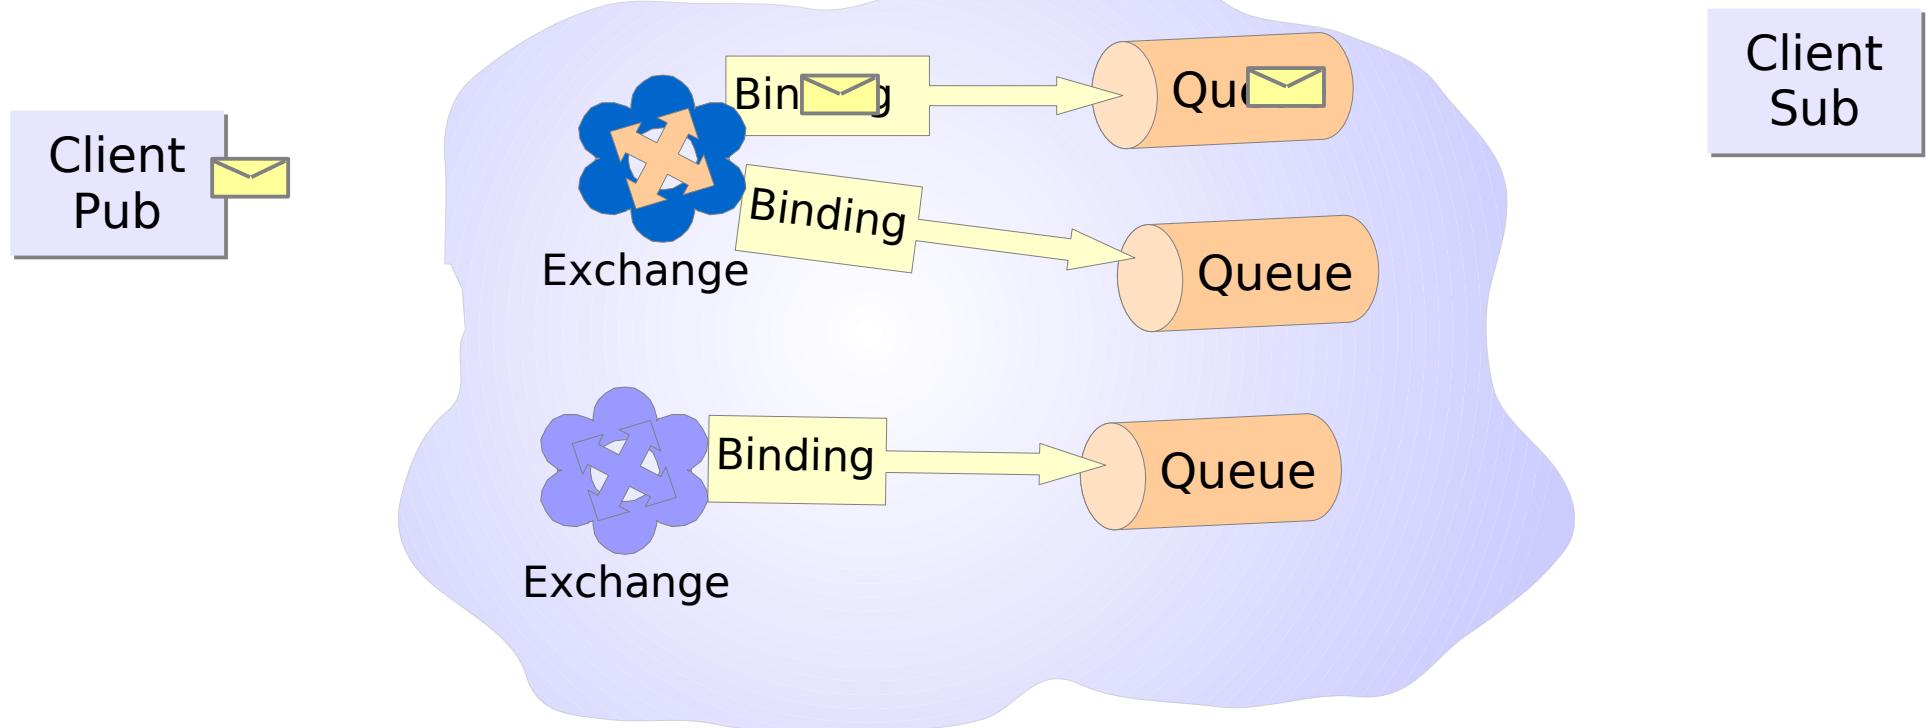
- JPMorgan
- Red Hat
- Deutsche Boerse
- Credit Suisse
- Goldman Sachs
- Cisco
- Iona
- Novell
- Microsoft
- Vmware
- ...

# AMQP Layers



# AMQP Model

## Shared Message Queue Space

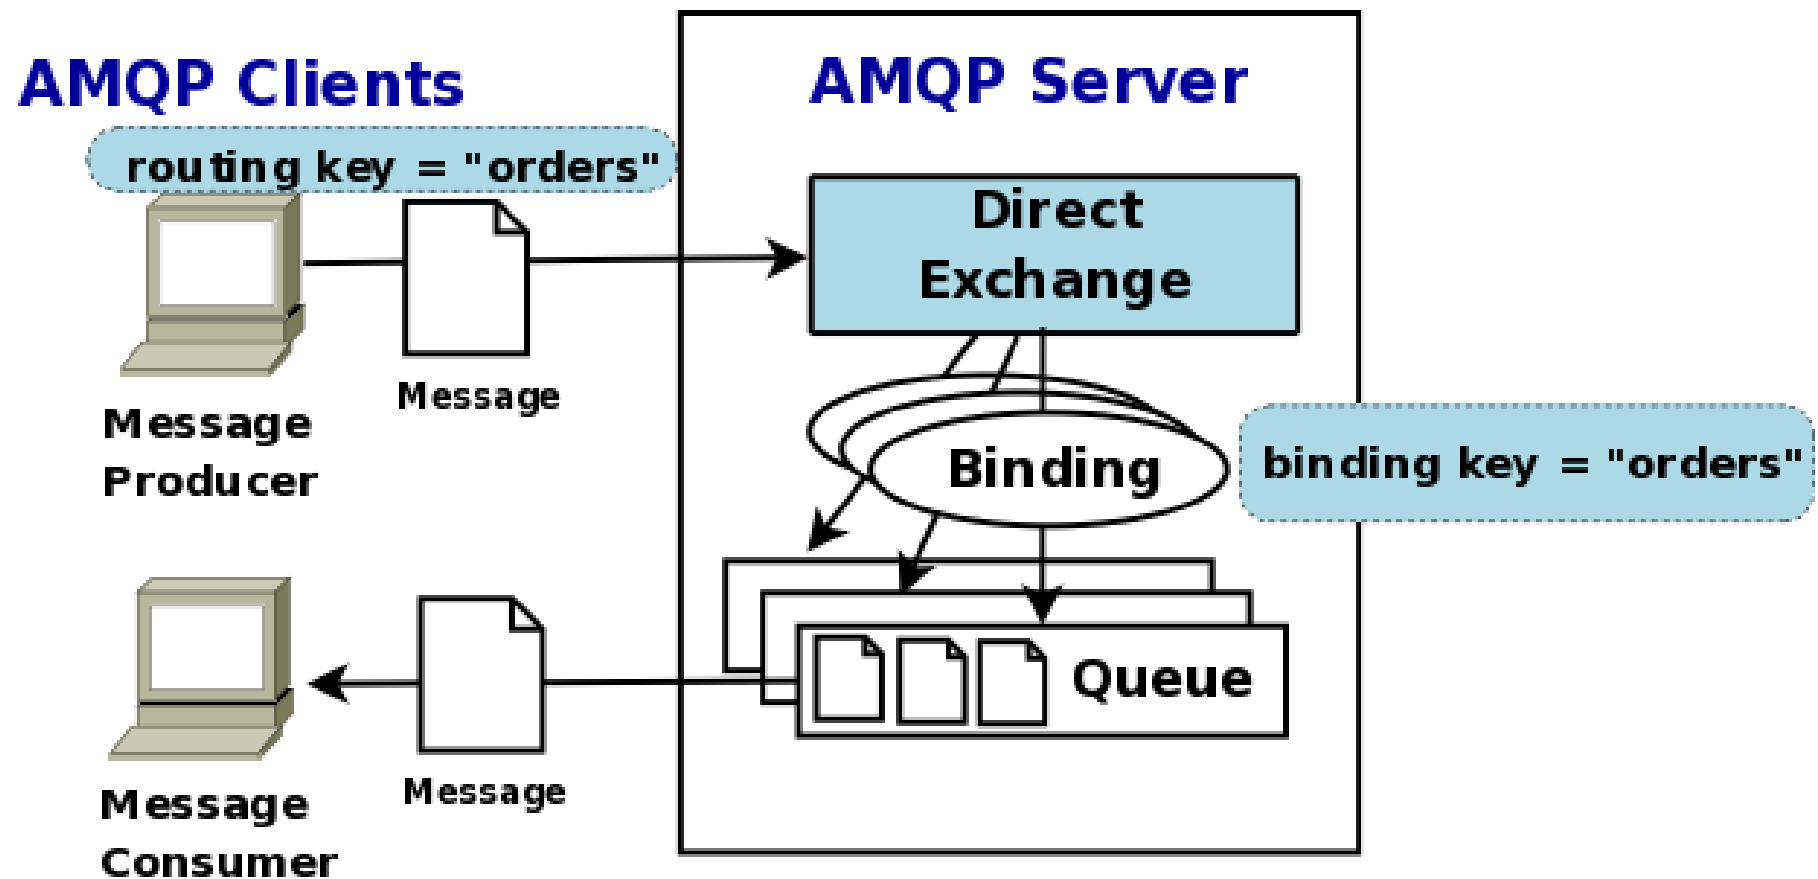


# AMQP Model

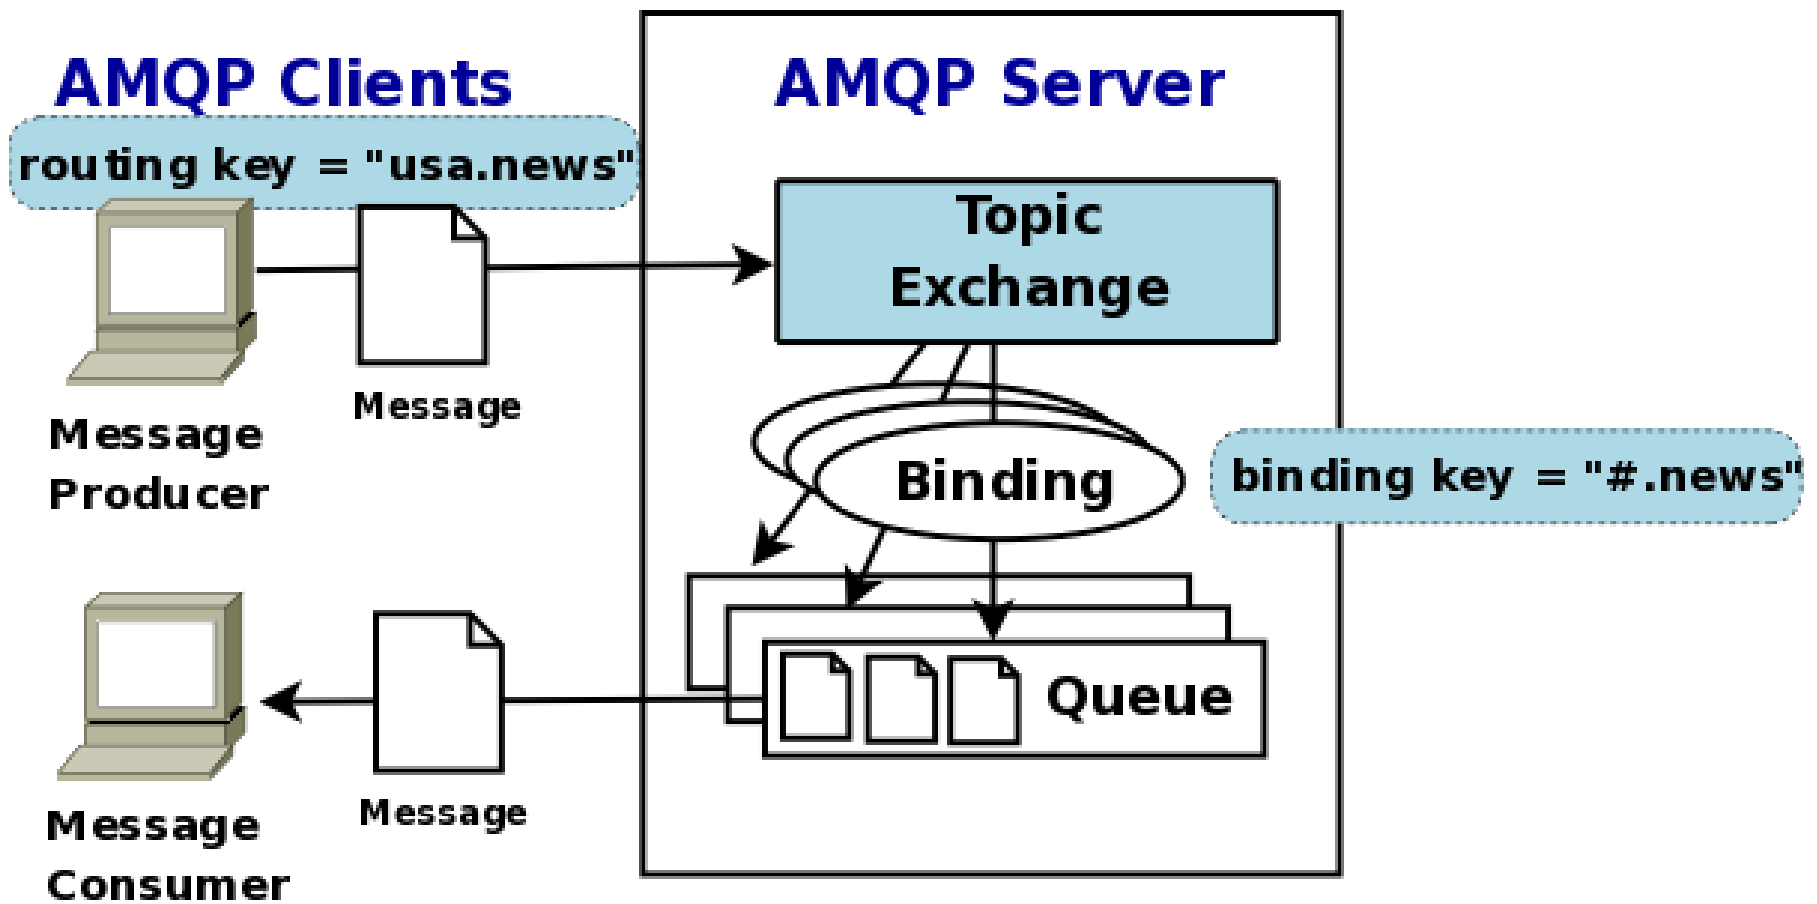
Provides a “Shared Queue Space” that is accessible to all interested applications.

- Message are published/sent to an **Exchange**
- Each message has an associated **Routing Key**
- Exchange forward messages to one or more **Message Queues** based on the **Routing Key**
- Consumers get messages from named **Message Queues**
- Only **one consumer** can get a given message from each Message Queue

# Direct Exchange



# Topic Exchange







# Exchange Types

- Direct
- Topic
- Fanout
- Headers
- XML
- Custom
- System

# JMS Model



# Mapping AMQP to JMS

- JMS Queue
  - Exchange type = Direct
  - One single queue
    - Routing key = Queue name = binding key
  - Queue is Public
    - All consumers consume from the same queue
  - Queue is Durable
    - Restored and Kept even if there isn't any consumer

# Mapping AMQP to JMS

- JMS Topic
  - Exchange type = Topic
  - One queue per consumer
    - Routing key = topic name
    - Binding key = wildcard
  - Queues are Private and
    - Volatile = standard subscriber
    - Durable = durable subscriber

## What about JMS Queue browsing ?

- JMS defines the notion of queue browsing
  - Messages are accessed but not consumed
- Can we push the concept even further?
  - The consumer should be able to decide whether
    - The message is of interest
    - Potentially consumed it

# Transfer of Responsibility

## ■ **No-acquire mode:**

- Only **data** is transferred, **NOT responsibility**
- No exclusive access to process the message
  - Another client may see, acquire and consume
- Need to explicitly acquire before processing

## ■ **Pre-acquire mode:**

- Both **data AND responsibility** are transferred
- Exclusive access to process the message
- No other client can see the message
- Can release to relinquish responsibility

## Accept Mode

### ■ **Explicit**

- Successful transfer is signaled by **semantic ack**

### ■ **None**

- Successful transfer upon **acquisition**

### ■ **Release**

- Relinquish responsibility for processing message that can be safely delivered to other clients

### ■ **Reject**

- Indicates a problem with processing a message (DLQ of the message)

# JMS Queue Browser

- JMS Queue browser
  - no-acquire - accept-mode = explicit  
send; look into message
  - More than JMS Queue browsing: consume messages of interest  
send; look into message; acquire; ack
- JMS Queue with guaranty delivery
  - pre-acquire - accept-mode = explicit
- Fast unreliable JMS Queue
  - pre-acquire - accept-mode = none



## Flow control issue

- **How can we handle fast producers?**
  - Messages are piling up on the server
  - Memory exhaustion
- **How can we speed up message delivery?**
  - Messages are sent one after each other
  - Increased latency
- **How can we handle small memory footprint?**
  - Cannot consume and or store big messages

# Message Flow

- Credit Based
  - Sender maintains credit balance with recipient
  - Credit Balance consist of a
    - Message Count
    - Byte Count
  - When a Message is sent both counts are decremented
  - When either value is zero no more messages are sent until further credit is received from peer.
  - If byte count is insufficient no partial messages can be sent

# Message Flow

## ■ Window

- Message acknowledgment implicitly grants:
  - a single unit of message credit
  - size of message in byte credits
- Controls the window of un-acked messages

# A Simple C++ Messaging Program

```
#include <qpid/client/Connection.h>
#include <qpid/client/Session.h>
#include <qpid/client/Message.h>

using namespace qpid::client;
using namespace qpid::framing;

int main() {
    const char* host = "127.0.0.1";
    int port = 5672;

    Connection connection;
    try {
        connection.open(host, port);
        Session session = connection.newSession();

        Message message;
        message.getDeliveryProperties().setRoutingKey("routing_key");
        message.setData("Hi, Mom!");

        session.messageTransfer(arg::content=message,
                               arg::destination="amq.direct");

        connection.close();
    } catch(const std::exception& error) {
        std::cout << error.what() << std::endl;
    }
}
```

# Using Java JMS with MRG Messaging

- Configure Queues, Exchanges – several possible ways
  - Using JNDI properties (see following slide)
  - Using the MRG Management Console
  - Using Low Level Java API
  - Using programs in C++ or Python
- Use vanilla Java JMS for messaging once configured

# Configuring Java JMS via JNDI Properties

- `connectionfactory.<jndiname>`  
The Connection URL used by the connection factory to create connections.
- `queue.<jndiname>`  
A JMS queue, implemented as an amq.direct exchange.
- `topic.<jndiname>`  
A JMS topic, which is implemented as an amq.topic exchange.
- `destination.<jndiname>`  
Can be used to define any amq destination, using a “Binding URL”.

# Example JNDI properties file for Java JMS

```
# JNDI properties file
java.naming.factory.initial =
org.apache.qpid.jndi.PropertiesFileInitialContextFactory
# register some connection factories
# connectionfactory.[jndiname] = [ConnectionURL]
# See MRG Messaging Tutorial for ConnectionURL format
connectionfactory.qpidConnectionFactory =
amqp://guest:guest@clientid/test?
brokerlist='tcp://localhost:5672'
# Register an AMQP destination in JNDI
# destination.[jndiName] = [BindingURL]
# See MRG Messaging Tutorial for BindingURL format
destination.directQueue =
direct://amq.direct//message_queue?routingkey='routing_key'
```

# Using JNDI to create Java JMS Session, Connection, Destination

```
// Load JNDI properties
Properties properties = new Properties();
properties.load(this.getClass().getResourceAsStream("direct.properties"));
// Create the JNDI initial context using JNDI properties
Context ctx = new InitialContext(properties);
// Look up Java JMS destination and connection factory
Destination destination = (Destination)ctx.lookup("directQueue");
ConnectionFactory connFact =
(ConnectionFactory)ctx.lookup("qpidConnectionFactory");
// Create Java JMS connection and the session using this
// connection factory
Connection connection = connFact.createConnection();
Session session = connection.createSession(false,
Session.AUTO_ACKNOWLEDGE);
```

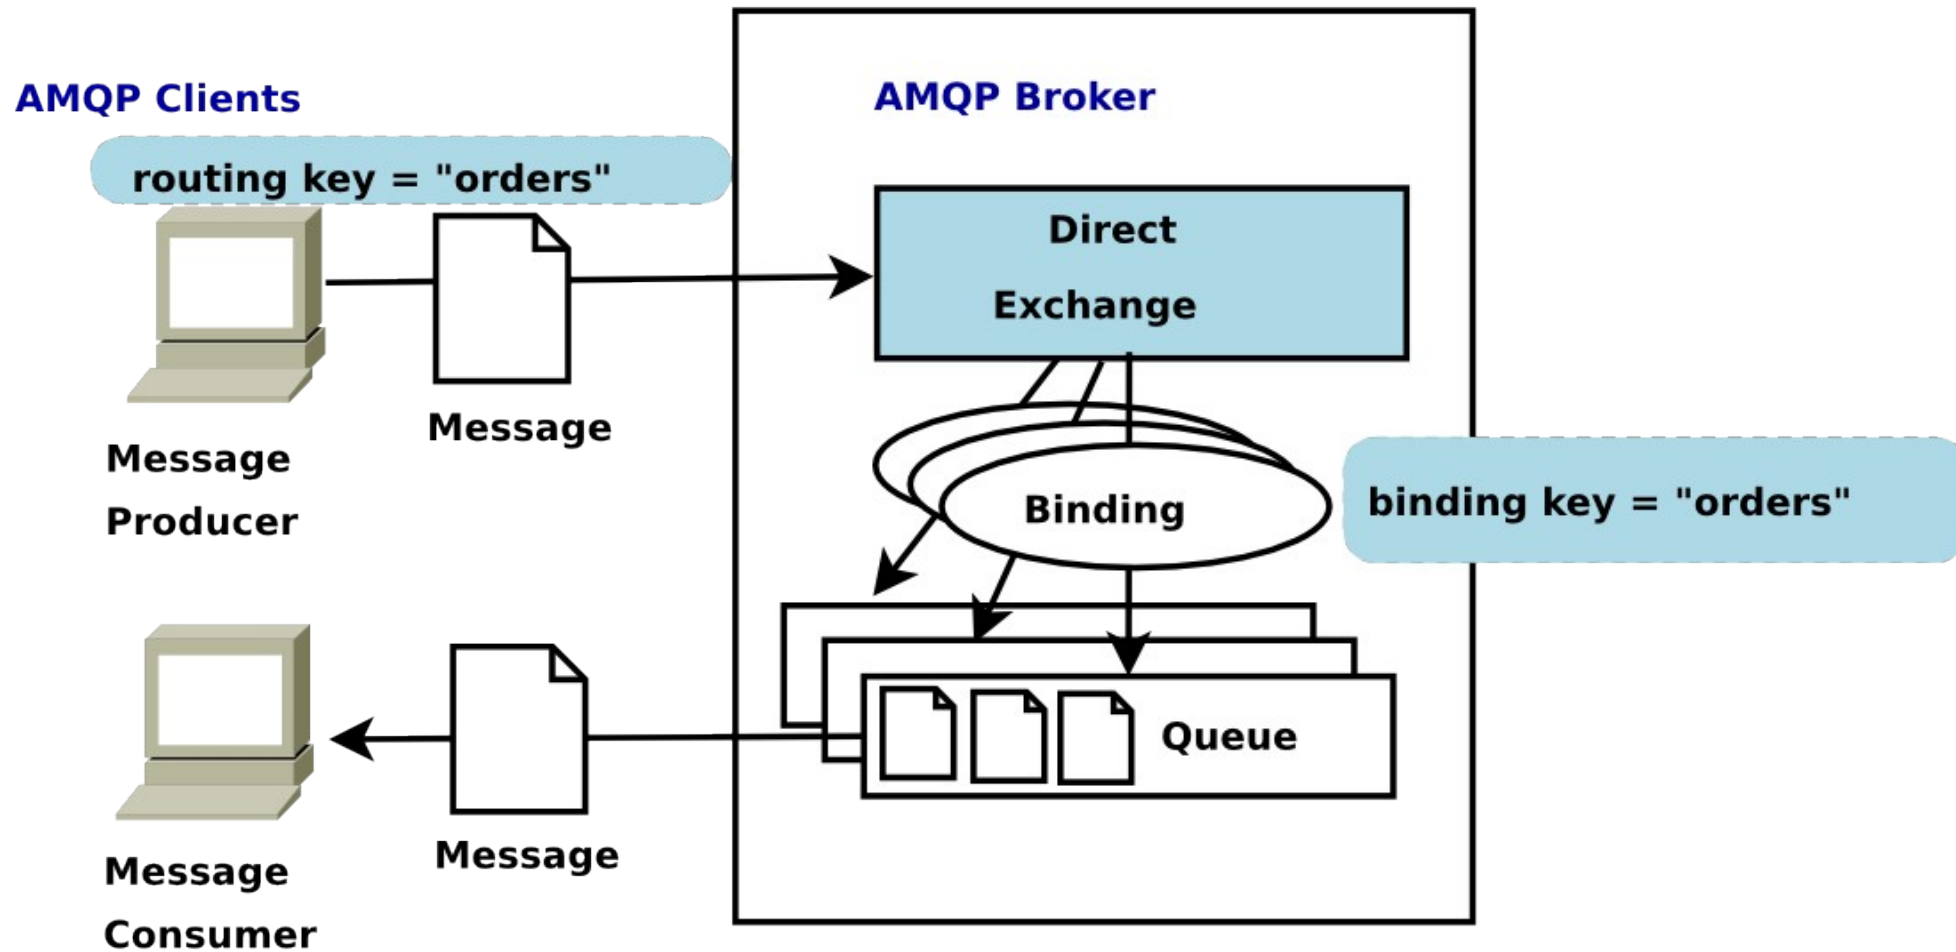


# Once configured, use the Java JMS API

**// Using standard Java JMS to send a message**

```
MessageProducer messageProducer =  
session.createProducer(destination);  
TextMessage message;  
message = session.createTextMessage("This is a text, this  
is only a text ...");  
messageProducer.send(message, getDeliveryMode(),  
    Message.DEFAULT_PRIORITY,  
    Message.DEFAULT_TIME_TO_LIVE);
```

# Direct Exchange: Point-to-Point



## Point-to-Point: Declaring a Queue and Binding

```
// arg::queue specifies the queue name
session.queueDeclare(arg::queue="message_queue");
// bind "message_queue" to "amq.direct" exchange
session.exchangeBind(arg::exchange="amq.direct",
    arg::queue="message_queue",
    arg::bindingKey="routing_key");
```

# Point-to-Point: Sending Messages

```
Message message;  
// Set routing key  
message.getDeliveryProperties().setRoutingKey("routing_key");  
// Send some messages  
for (int i=0; i<10; i++) {  
    stringstream message_data;  
    message_data << "Message " << i;  
    message.setData(message_data.str());  
    session.messageTransfer(arg::content=message,  
                           arg::destination="amq.direct");  
}  
message.setData("That's all, folks!");  
session.messageTransfer(arg::content=message,  
                       arg::destination="amq.direct");
```

# Point-to-Point: Receiving Messages

**// Create a Listener, Derived from MessageListener**

```
class Listener : public MessageListener {  
    private:  
        SubscriptionManager& subscriptions;
```

```
    public:
```

```
        Listener(SubscriptionManager& subscriptions);  
        virtual void received(Message& message);
```

```
};
```

**// Implement constructor**

```
Listener::Listener(SubscriptionManager& subs) : subscriptions(subs) {}
```

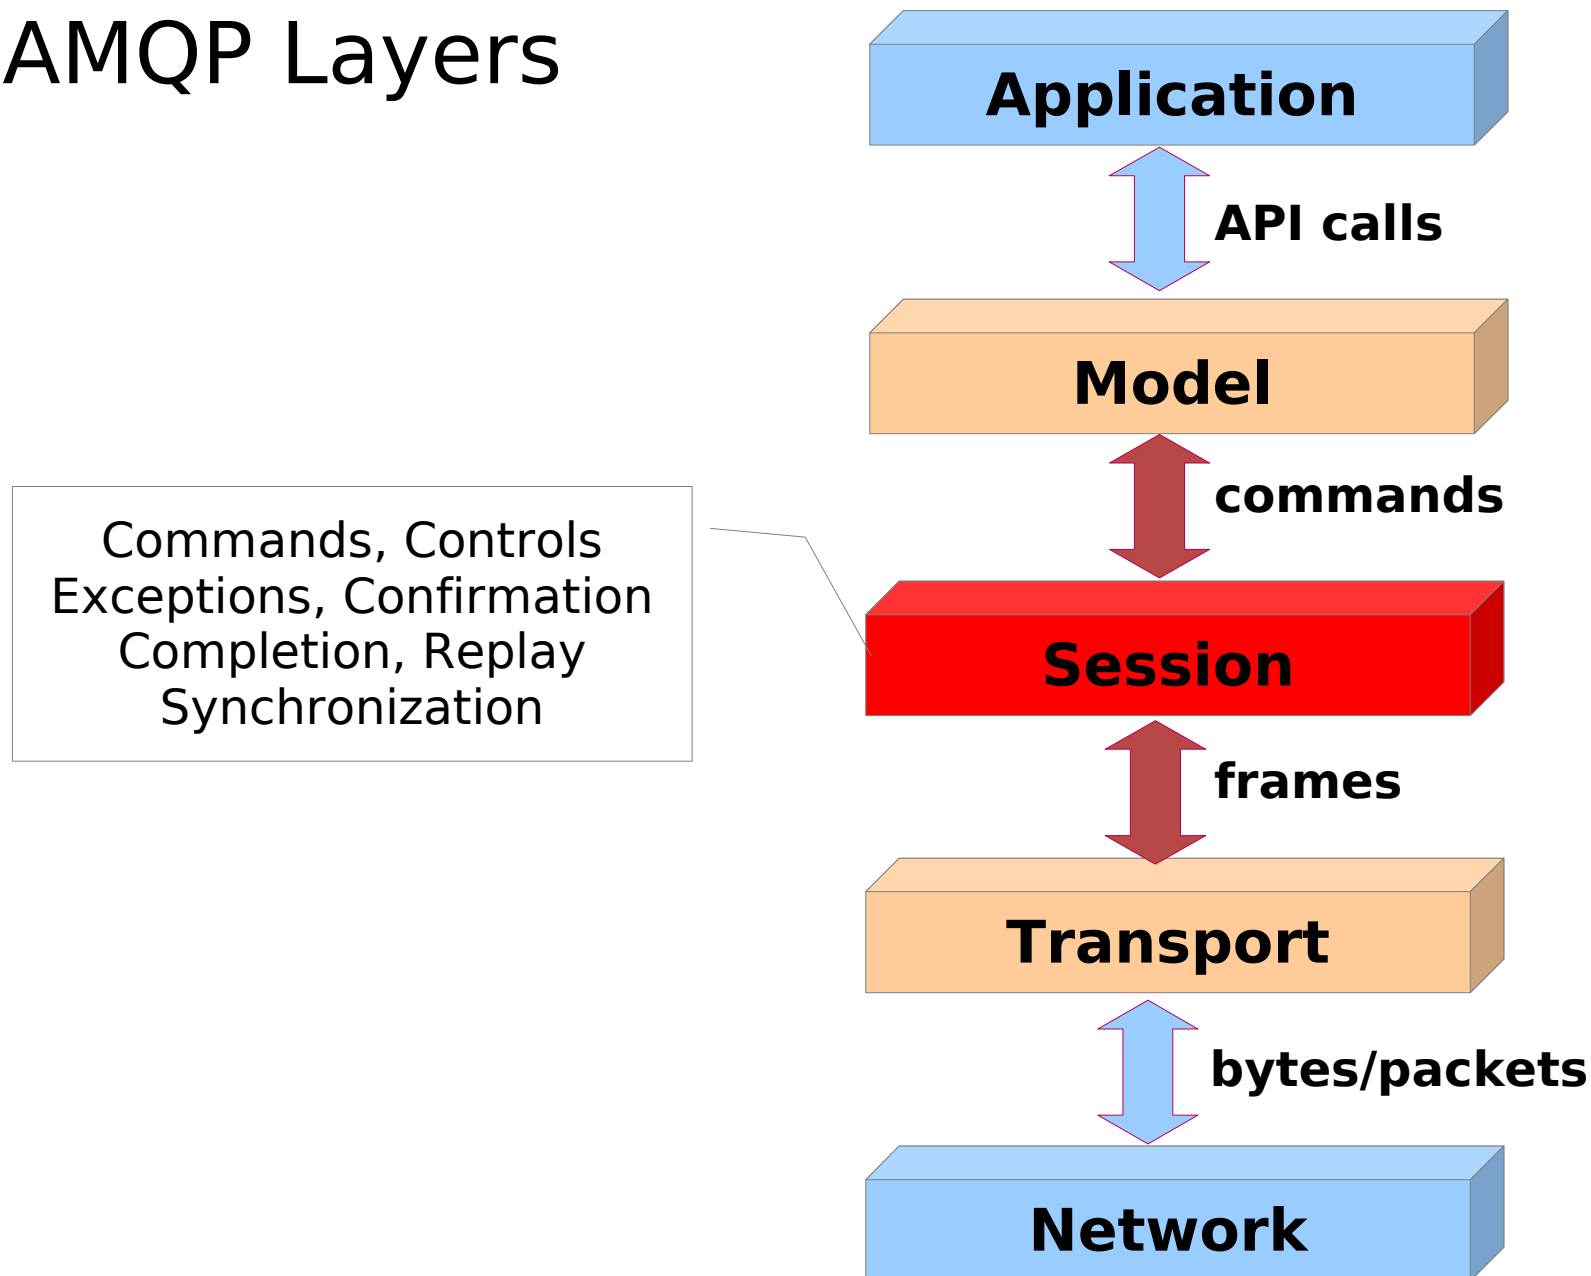
**// Implement Listener::received()**

```
void Listener::received(Message& message) {  
    std::cout << "Message: " << message.getData() << std::endl;  
    if (message.getData() == "That's all, folks!") {  
        std::cout << "Shutting down listener for "  
                    << message.getDestination()  
                    << std::endl;  
        subscriptions.cancel(message.getDestination());  
    }  
}
```

# Point-to-Point: Receiving Messages

```
// Subscribe Listener to "message_queue"  
SubscriptionManager subscriptions(session);  
Listener listener(subscriptions);  
subscriptions.subscribe(listener, "message_queue");  
// Receive messages until Listener::received() cancels  
subscription  
subscriptions.run();
```

# AMQP Layers



## Reliability issues

- How can we handle network failures?
  - Message transfer can be interrupted
  - Clients needs to reconnect
- Example

A train sends messages to the Railway company HQ

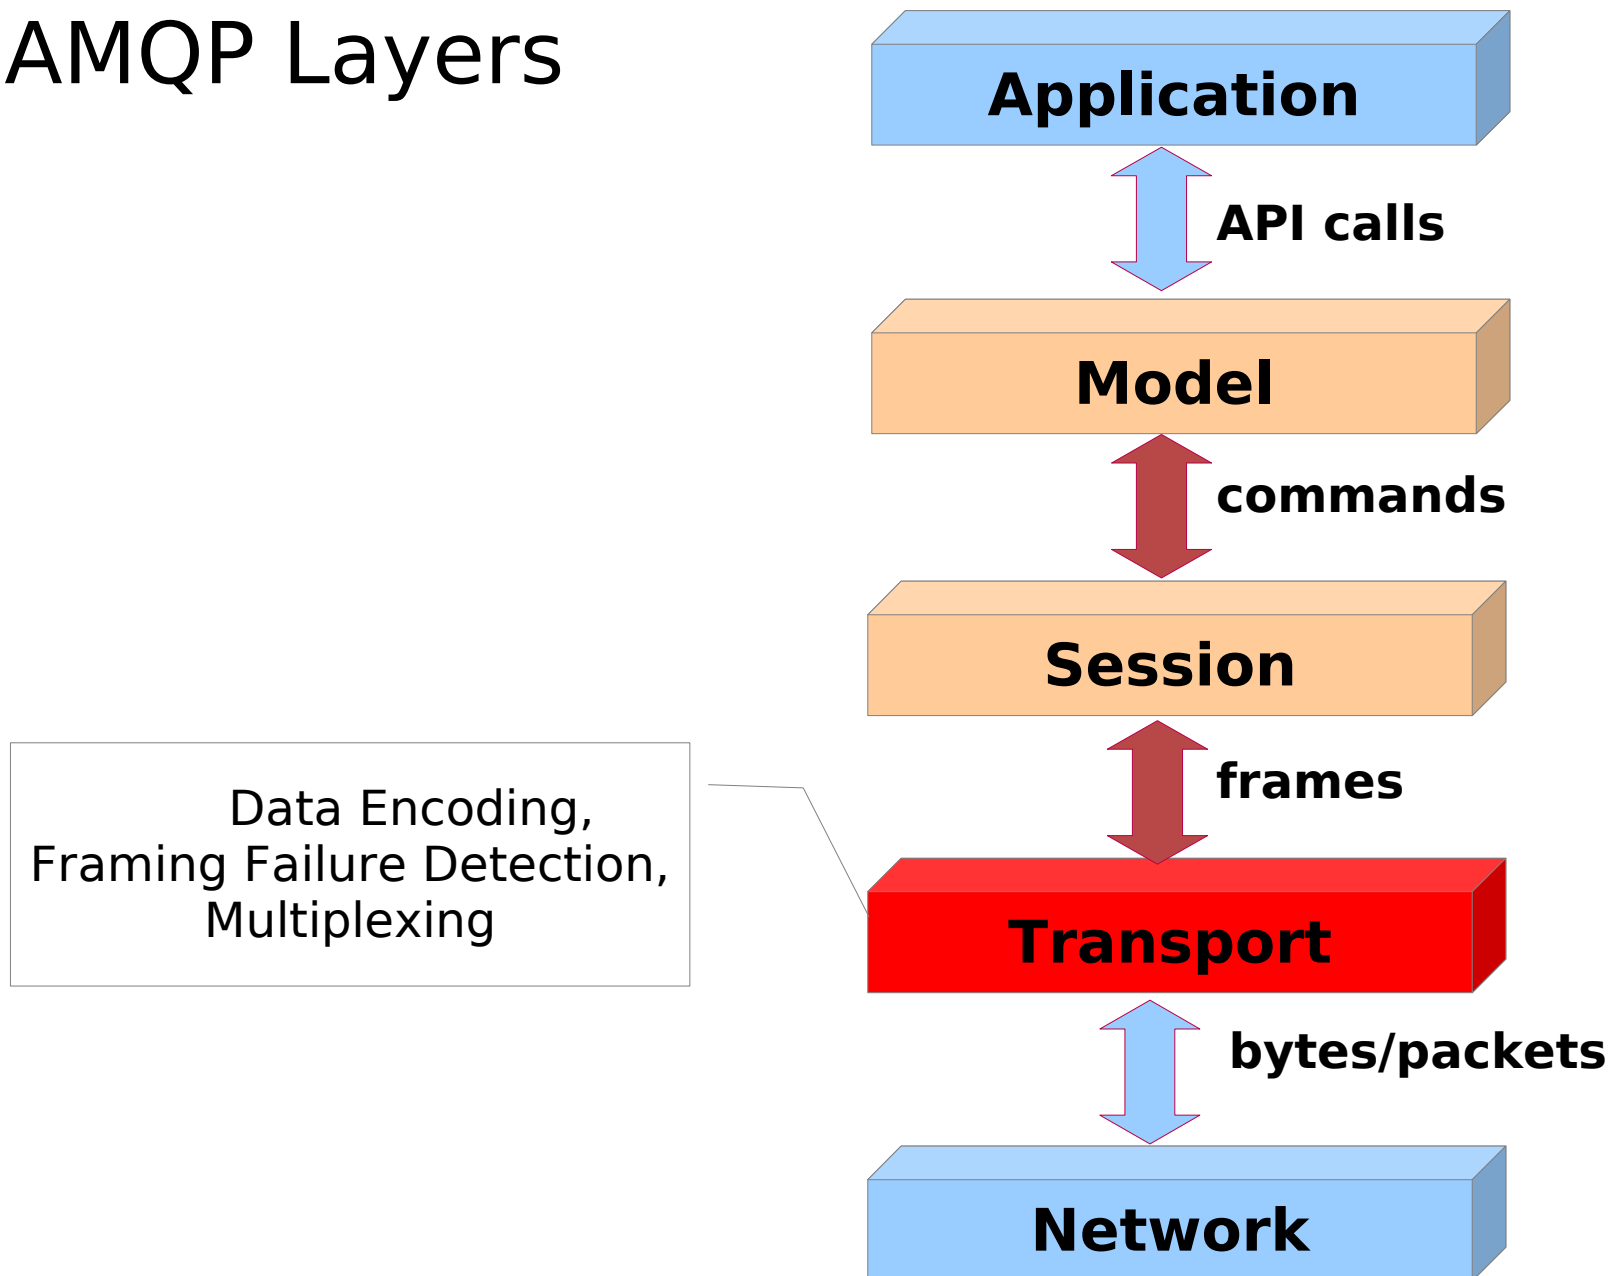
  - The train is moving!
  - The network nature is changing
    - Satellite, GSM, Etc.
- In cluster mode how can we transparently switch nodes



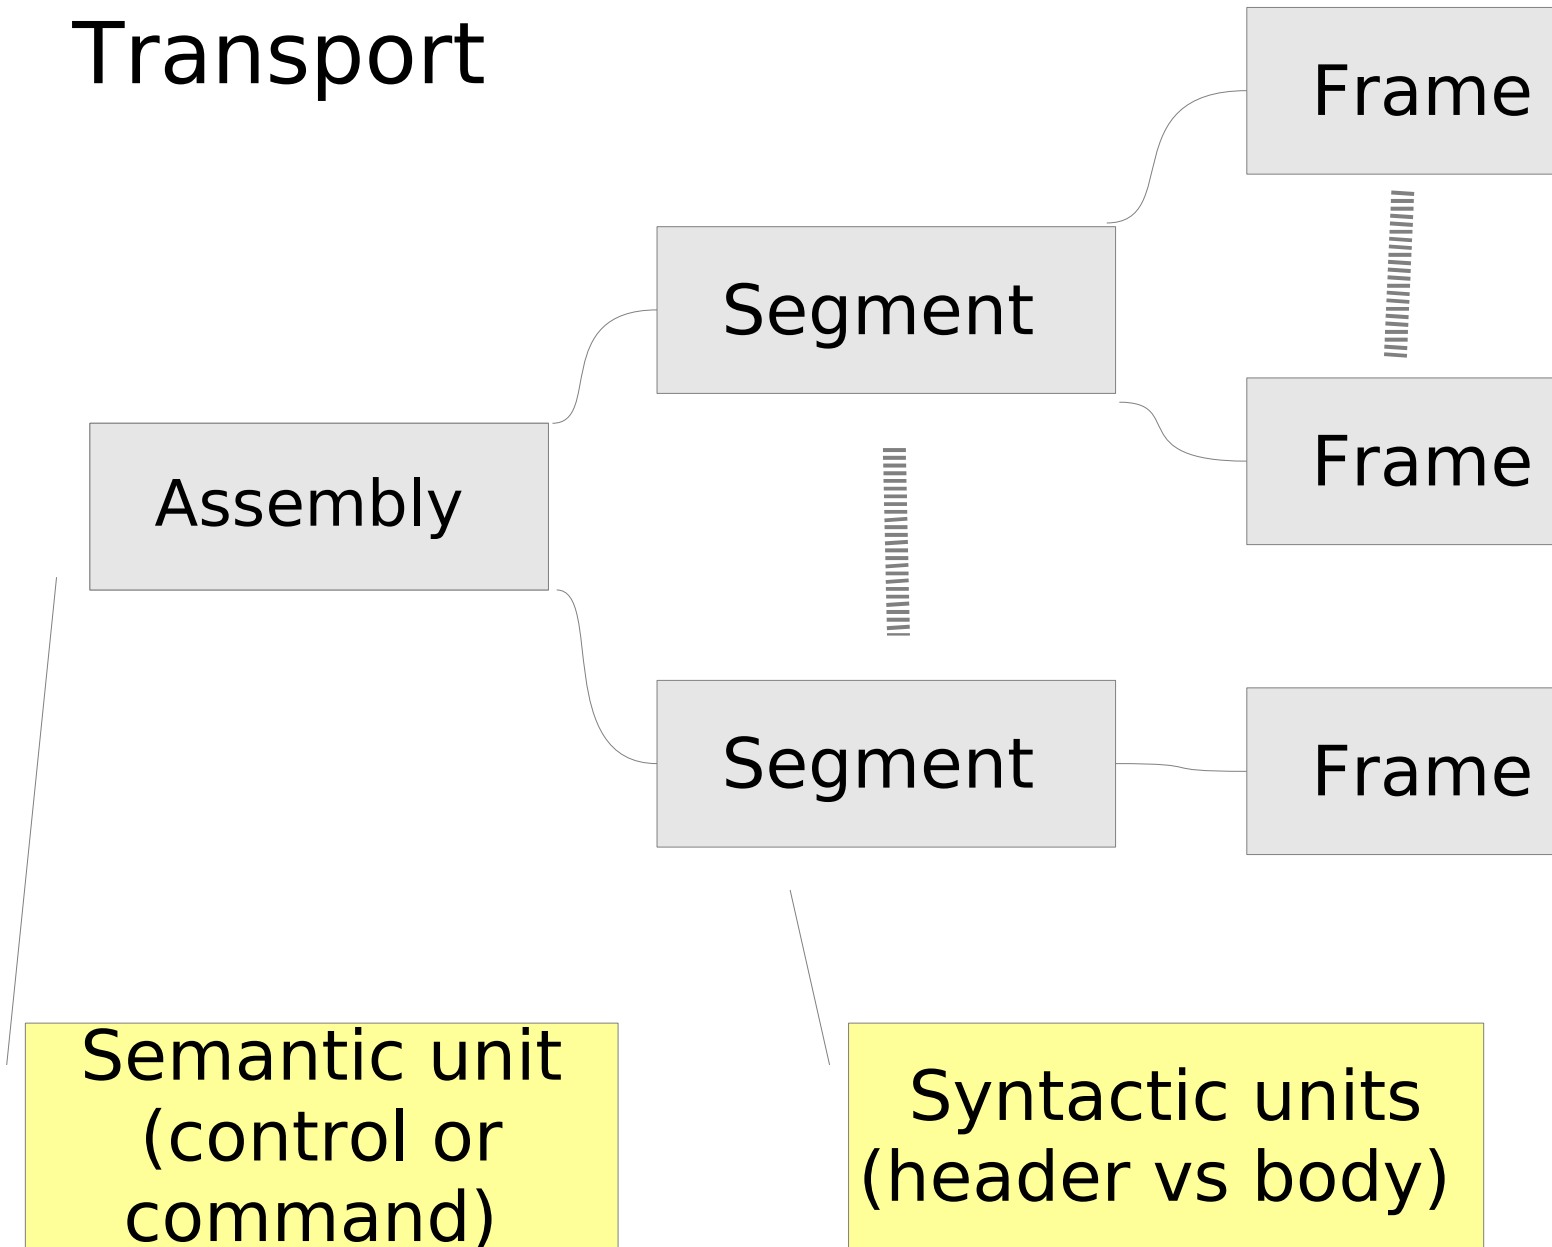
# Session

- Sequential numbering scheme with rollover to identify each **command uniquely** within a session
- Session state
  - a **replay buffer** of full or partial commands which a peer does not yet have confirmation its partner has received
  - an **idempotency barrier** - a set of commands identifier which the peer knows that it has received but cannot be sure that its partner will not attempt to re-send.

# AMQP Layers




# Transport



# Frames

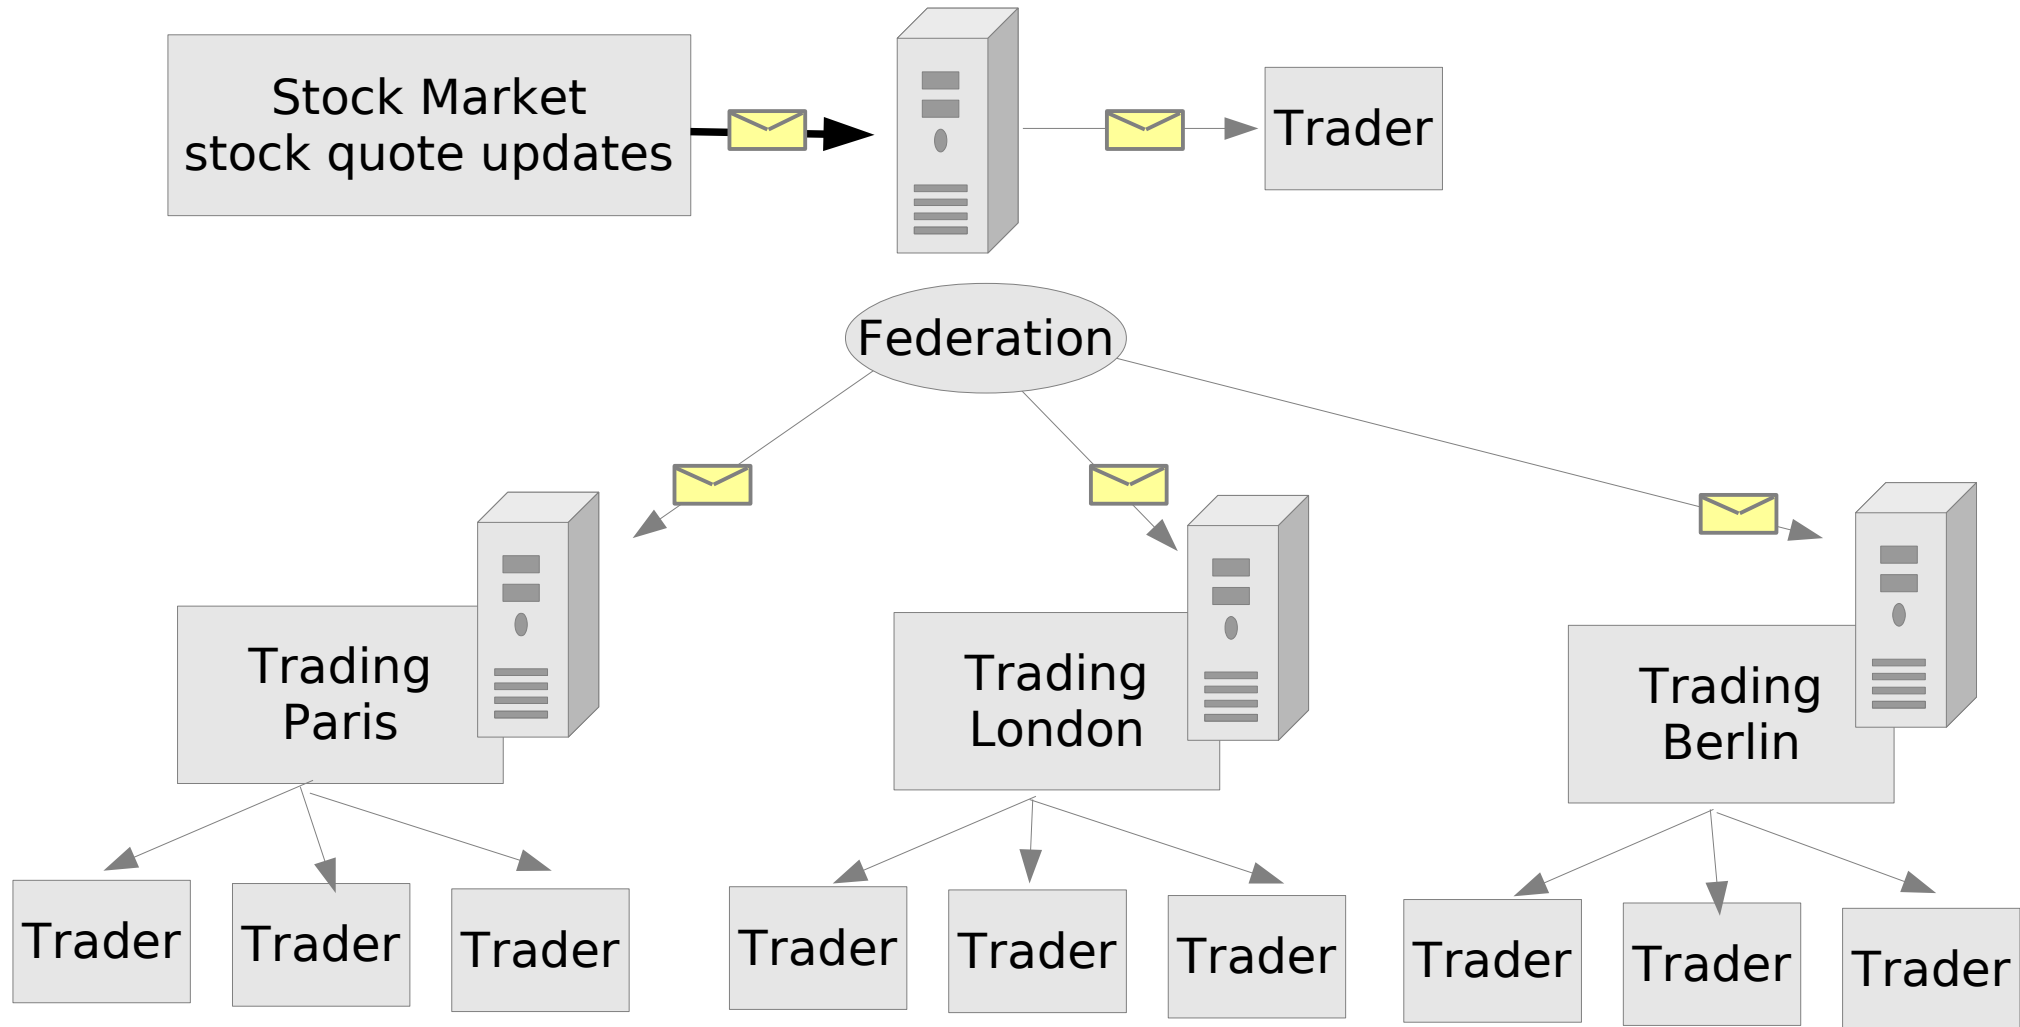
- Header:
  - Channel
    - divides a single frame transport into distinct channels (sessions)



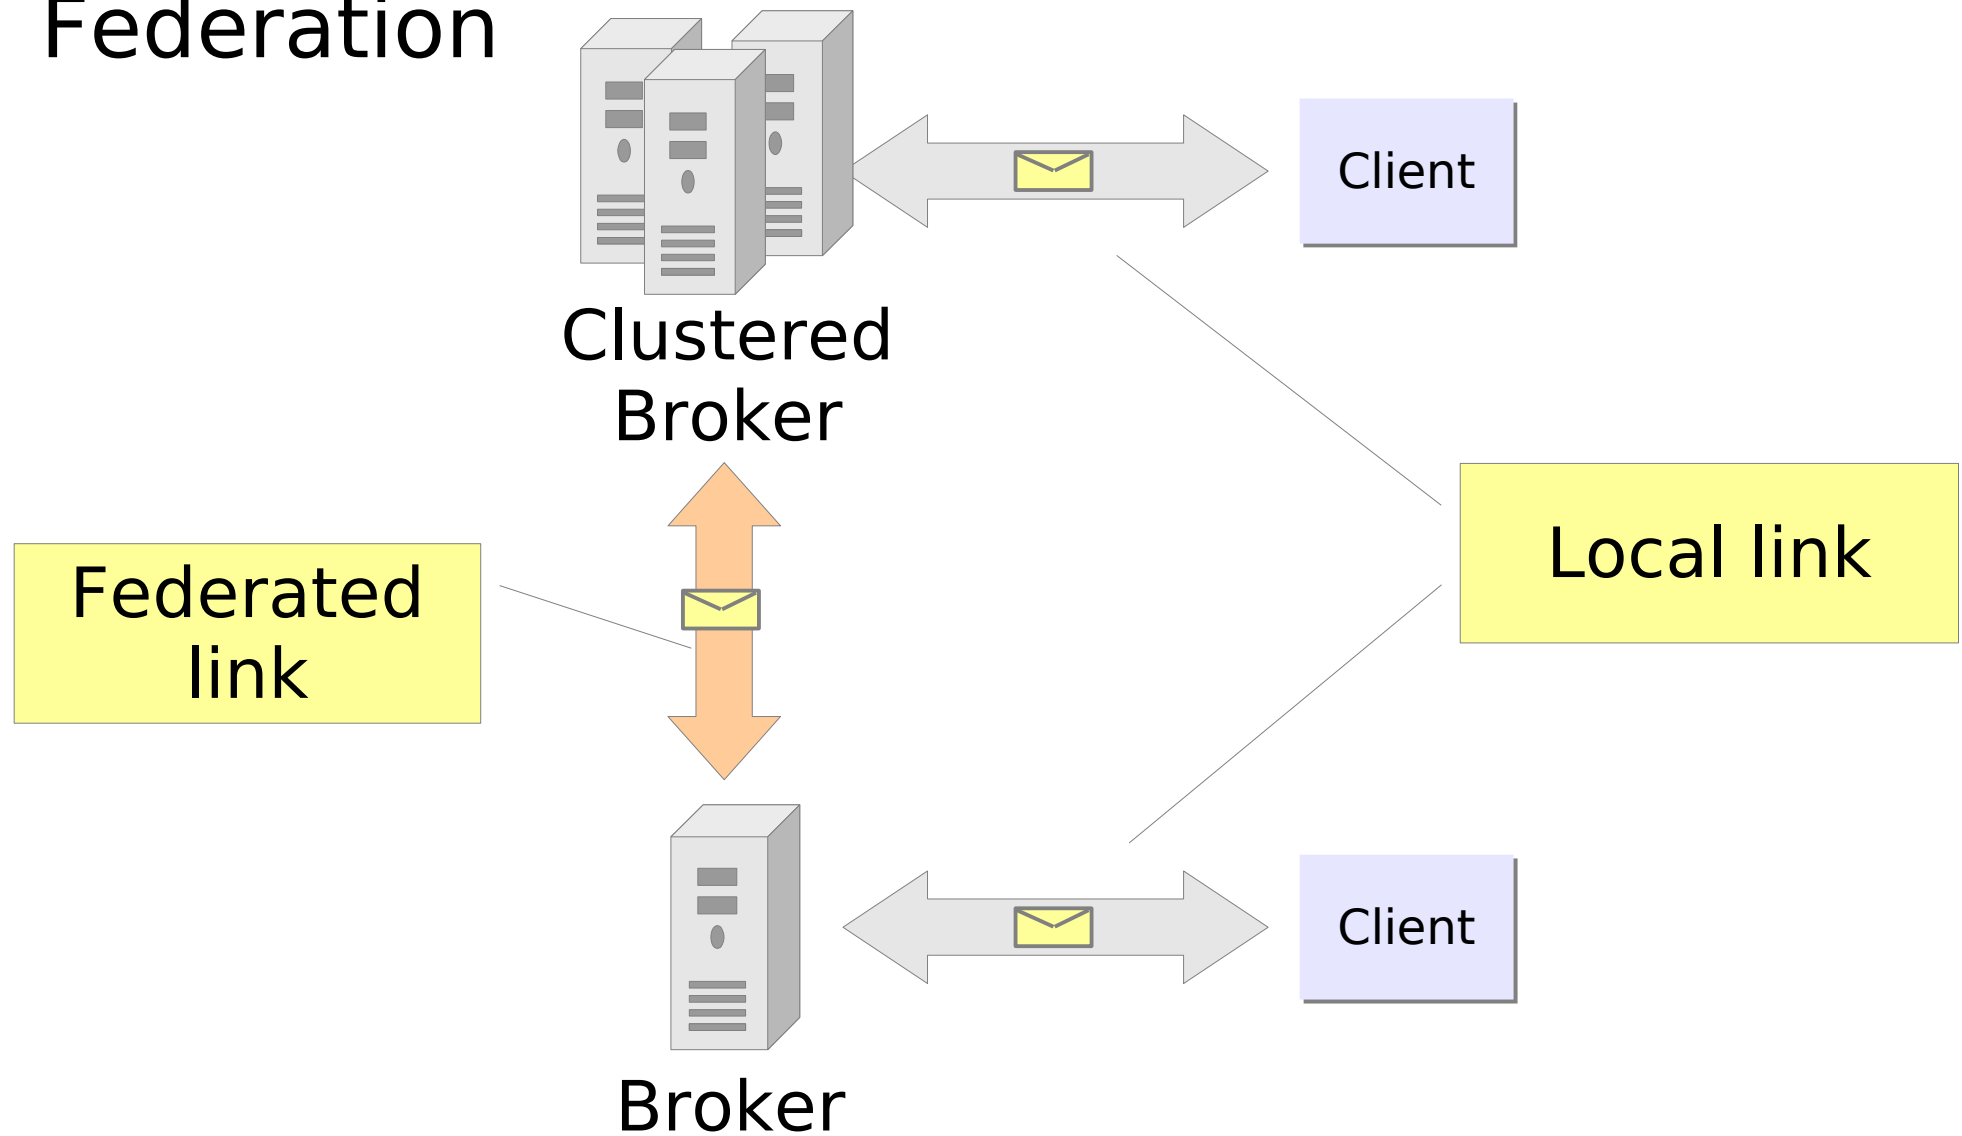
# Federation

- Federation
  - Joining of multiple brokers together in a large functioning network
- Clustering
  - Several brokers deployed that act as a single broker
    - High availability
    - Performance

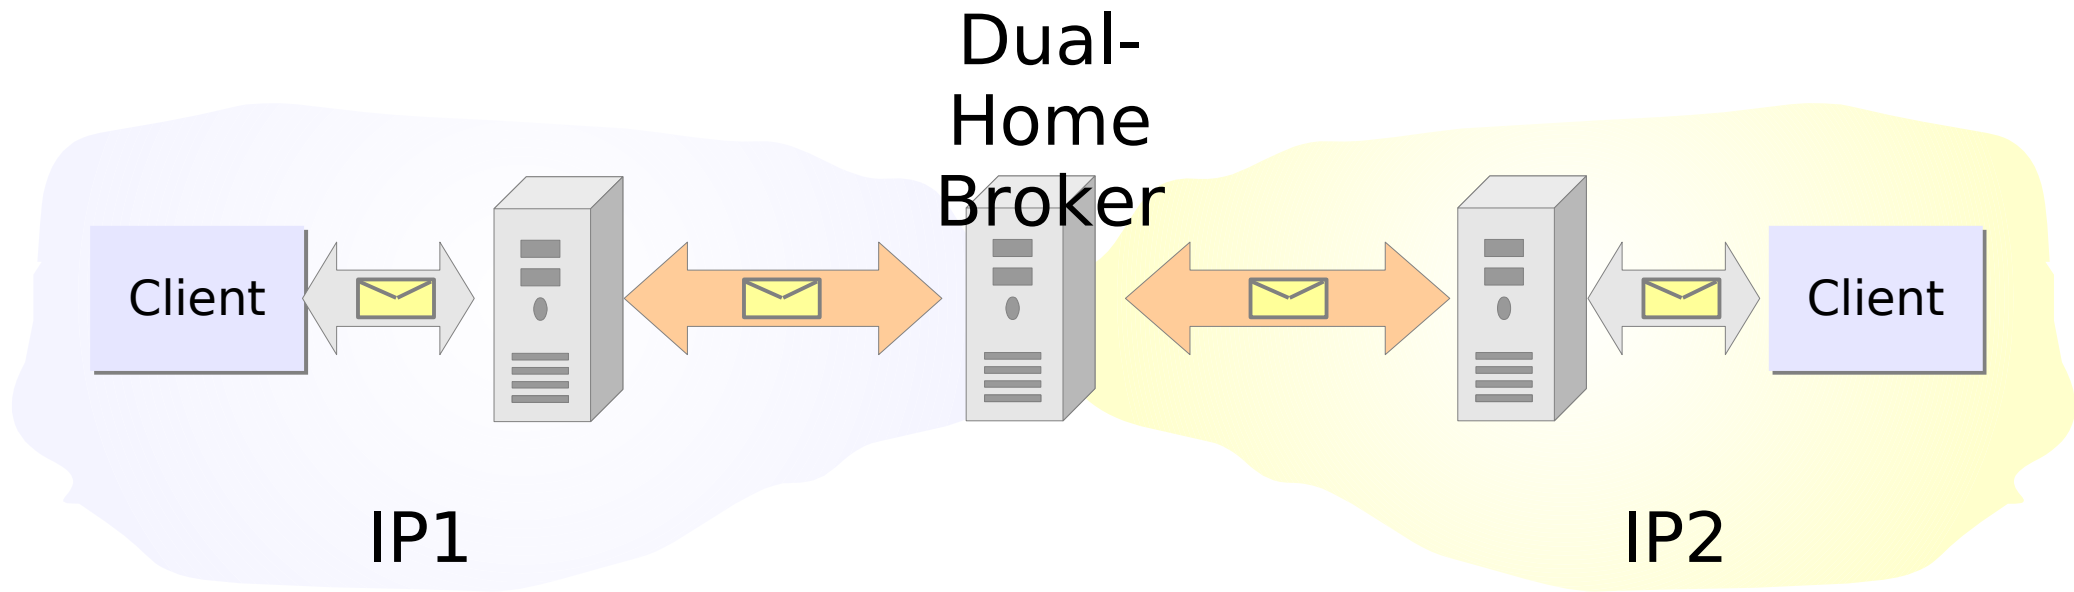
# Federation



# Federation

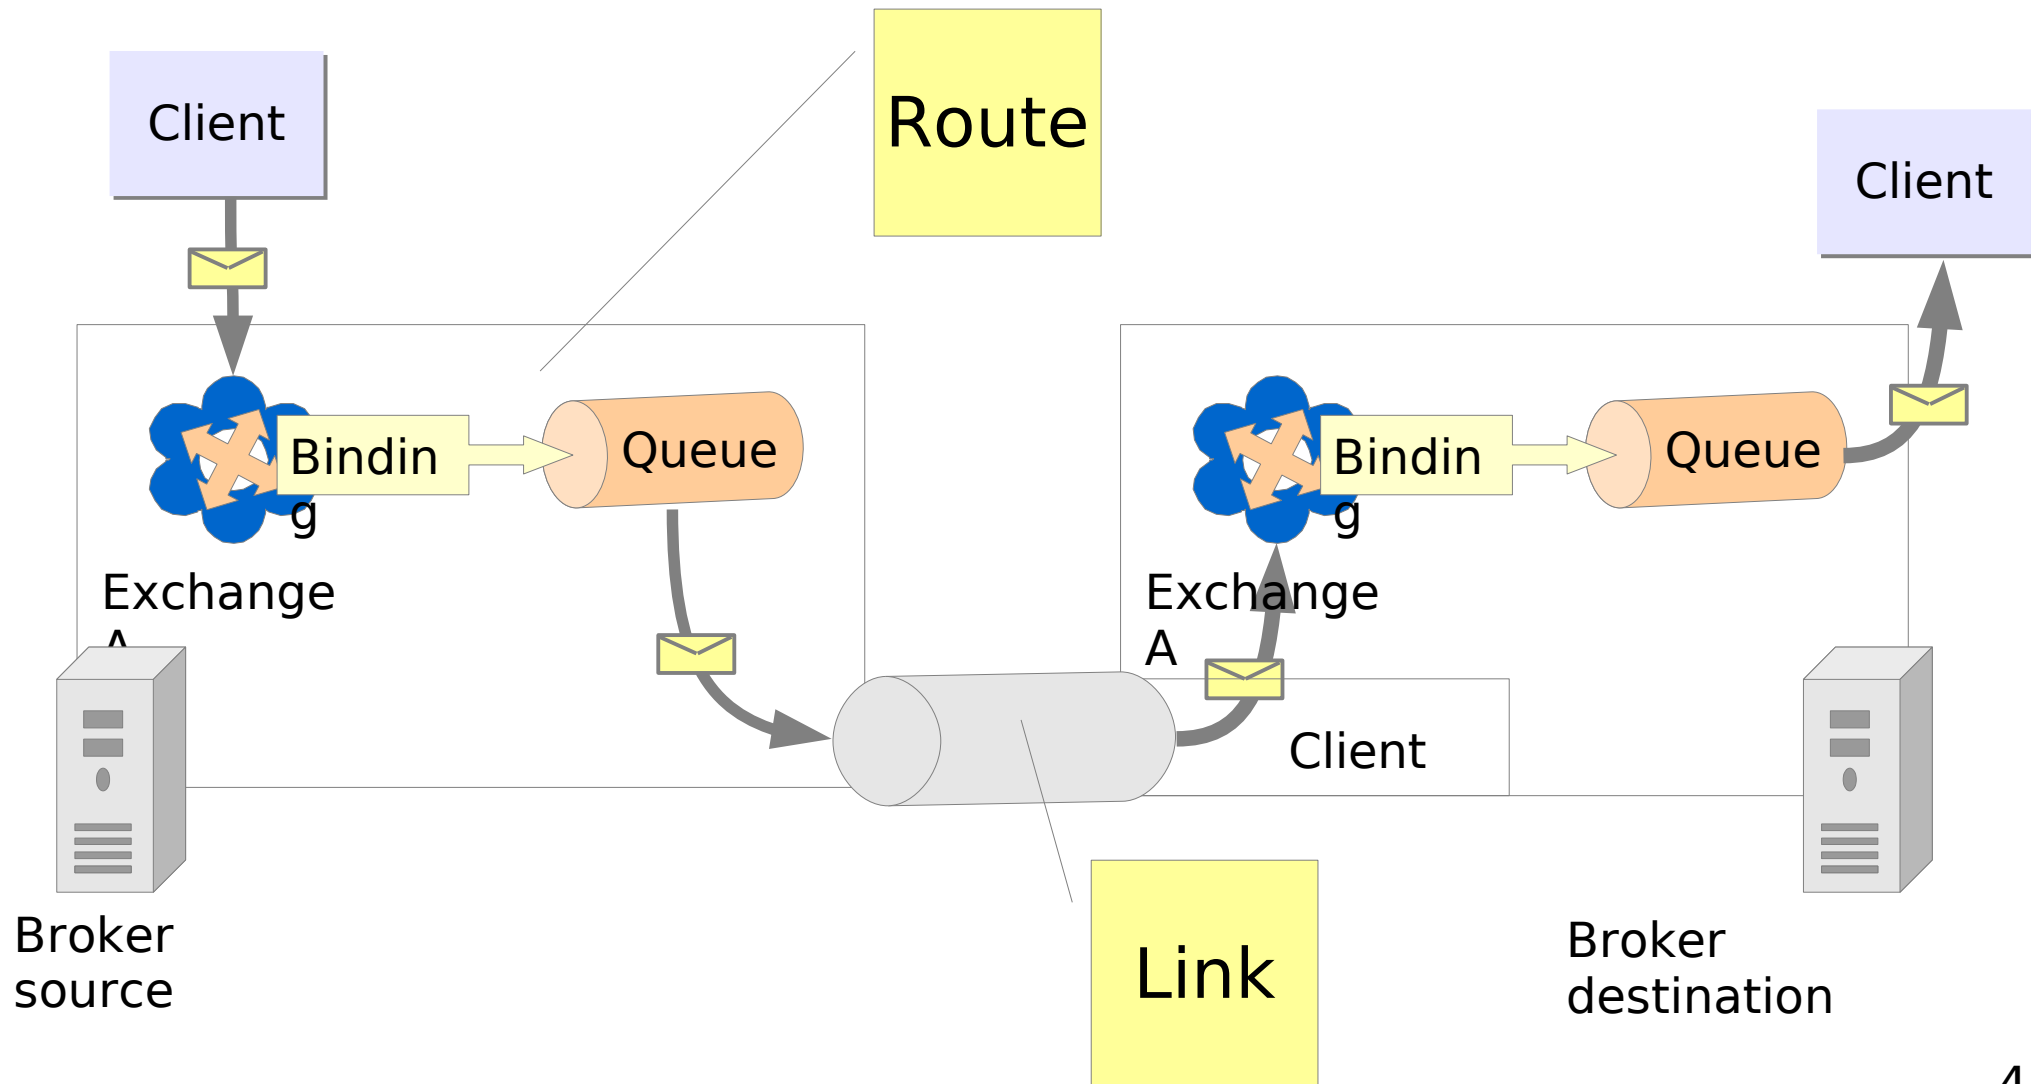


# Isolated Networks






# Distributed Exchange



# Federation


- Link:
  - Source broker
  - Destination broker
- Route:
  - Source broker
  - Destination broker
  - Exchange name
  - Binding key

# Federation - cycles



# Message Store

- Asynchronous Journal
  - Asynchronous IO
  - O-direct flag (disable buffering)
  - “Circular buffer”



# Authentication SASL

- Simple Authentication and Security Layer
  - Decouples authentication mechanisms from AMQP
  - Supports a rich set of mechanisms:
    - Anonymous, Plain, MD5 challenge/response
    - NTLM for Windows
    - GSSAPI for Kerberos v5
    - EXTERNAL for x.509 Certificate authentication
- Authentication of AMQP client

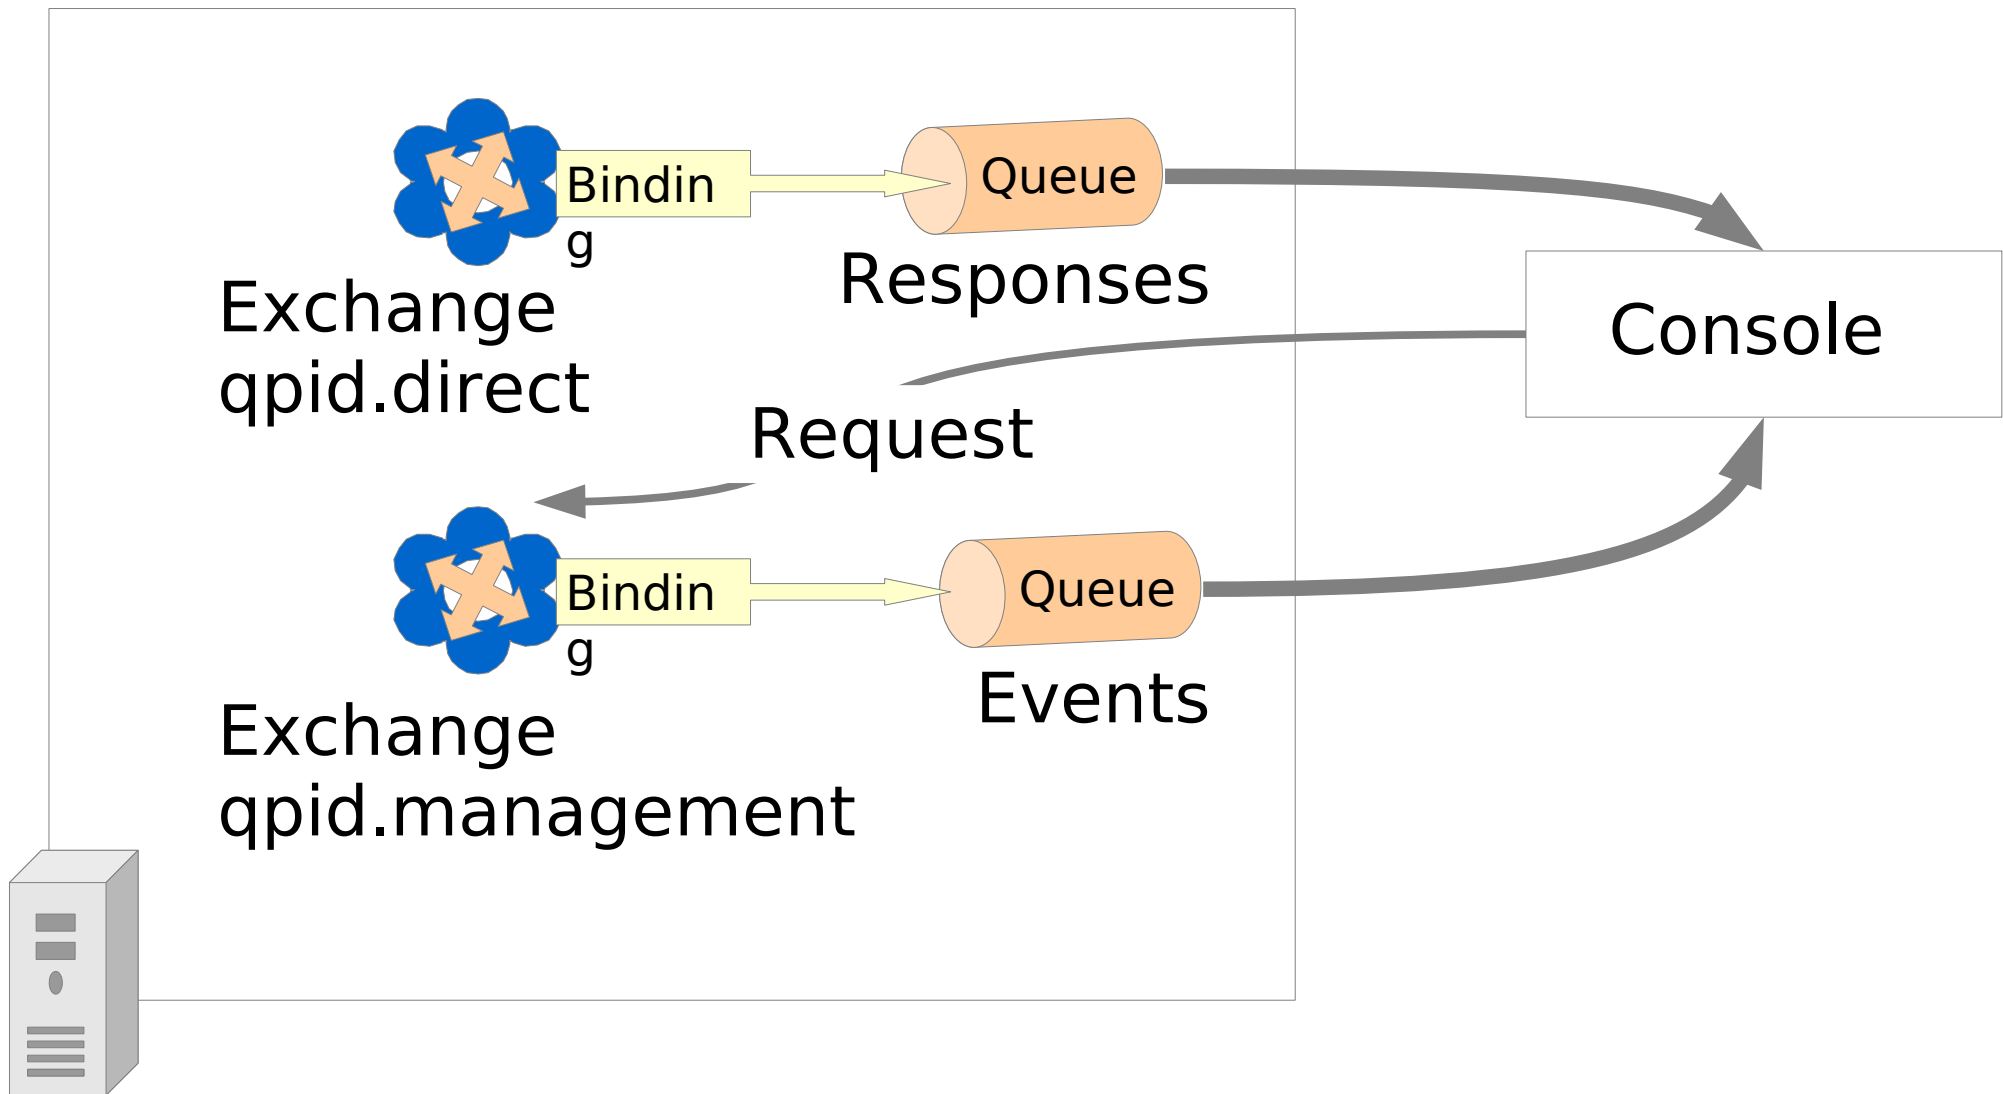
# Message-level Security - Signing

- Protects data integrity end-to-end
- Signature covers Header and Body segments concatenated
- AMQP 1.0
  - Add a new footer segment to contain the signature of the concatenation of the header and body segments
  - Add a new field to message-properties: signature-control. This is used to identify the signing mechanism used for the message.
  - Add reject codes were added to allow a consumer to reject a message (invalid-signature, missing-signature, untrusted-signature)

## Message-level Security – Encryption

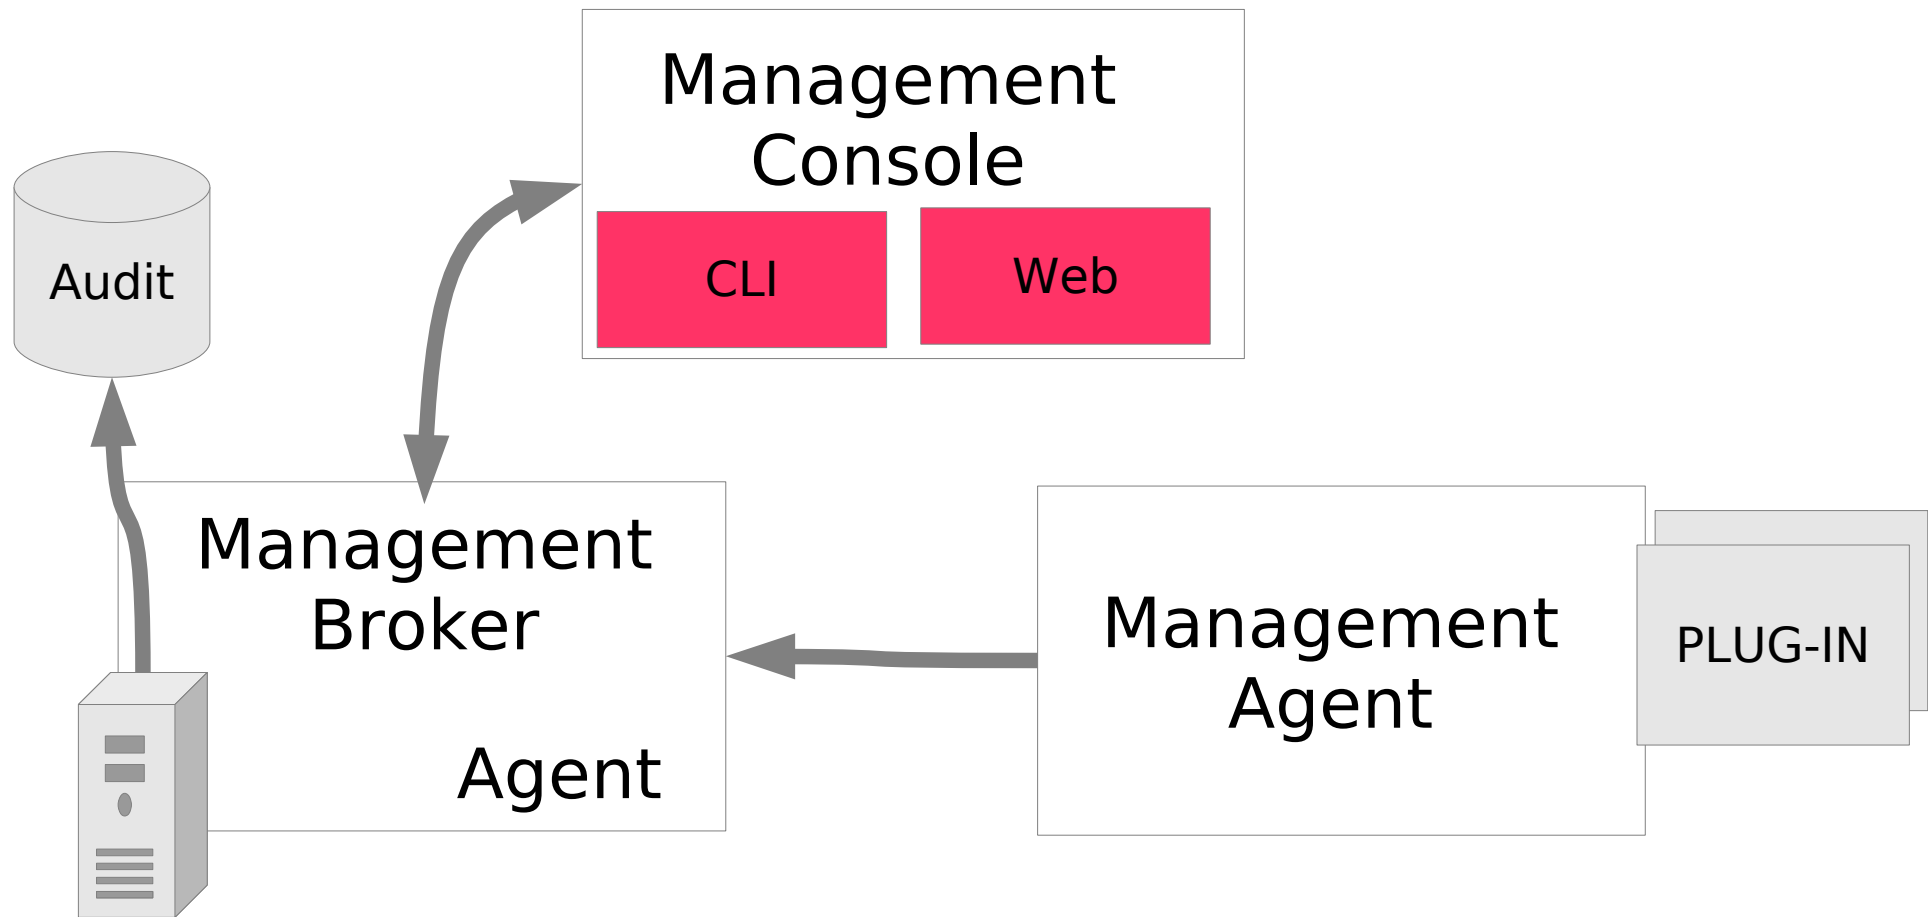
- Encryption may be applied to the entire body segment
  - Broker need not be trusted
- The header segment must not be encrypted:
  - Header contains information needed by the broker
  - Header contains information needed to decrypt the body
- AMQP SP1
  - Add field: encryption-control – used by consumer to identify the algorithm and key to decrypt the body (S-MIME, other mechanisms)
- **Key exchange** is outside the scope of AMQP

# Management





# Management



# Questions?



## Sorting out the terms

Term	Who	What
AMQP	AMQP Working Group	Advanced Message Queuing Protocol
Qpid	ASF	Implementation of AMQP
BdbStore	Red Hat	RedHat's durable message store (QPID plug-in)
MRG	Red Hat	Product: Messaging, Real Time, Grid